

## HIVE CODES:

Create database db

Show databases present in hive

Use database db to work on

```
hive> create database if not exists db;
OK
Time taken: 0.227 seconds
hive> show databases;
OK
bda01
db
default
Time taken: 0.062 seconds, Fetched: 3 row(s)
hive> use db;
OK
Time taken: 0.044 seconds
hive>
```

Create an internal table student with the column names student id, student name ,marks ,department number

```
hive> create table if not exists student(sid int,sname string,marks int,dpno int) row format delimited fields terminated by ',';
OK
Time taken: 0.572 seconds
hive>
```

Structure of the table

```
hive> describe student;
OK
sid                int
sname              string
marks              int
dpno               int
Time taken: 0.263 seconds, Fetched: 4 row(s)
hive>
```

Inserting the data from local disk

```
hive> load data local inpath "/home/vboxuser/hadoopdata/data.csv" into table student;
Loading data to table db.student
OK
Time taken: 0.624 seconds
hive>
```

```
vboxuser@vanditha:~$ cd hadoopdata
vboxuser@vanditha:~/hadoopdata$ nano data.csv
vboxuser@vanditha:~/hadoopdata$ pwd
/home/vboxuser/hadoopdata
vboxuser@vanditha:~/hadoopdata$
```

Displays the content of the table with column names

```
Time taken: 0.624 seconds
hive> SET hive.cli.print.header=true;
hive> select * from student;
OK
student.sid      student.sname    student.marks    student.djno
1                raj              80               12
2                ram              81               11
3                sam              88               12
4                rita             70               10
NULL            NULL             NULL             NULL
Time taken: 2.165 seconds, Fetched: 5 row(s)
hive>
```

Create external table

```
Time taken: 2.165 seconds, Fetched: 5 row(s)
hive> create external table if not exists student1(sid int,sname string,marks int,djno int) row format delimited fields terminated by ',';
OK
Time taken: 0.086 seconds
hive> describe student1;
OK
col_name      data_type    comment
sid           int
sname         string
marks         int
djno          int
Time taken: 0.076 seconds, Fetched: 4 row(s)
```

Moving the data from local disk to hdfs

```
vboxuser@vanditha:~/hadoopdata$ hdfs dfs -ls /user/data
Found 1 items
-rw-r--r--  1 vboxuser supergroup      72 2023-04-21 13:59 /user/data/empdata_pig.csv
vboxuser@vanditha:~/hadoopdata$ hdfs dfs -put "/home/vboxuser/hadoopdata/data.csv" "/user/data/data.csv"
vboxuser@vanditha:~/hadoopdata$ hdfs dfs -ls /user/data
Found 2 items
-rw-r--r--  1 vboxuser supergroup      50 2023-04-28 23:58 /user/data/data.csv
-rw-r--r--  1 vboxuser supergroup      72 2023-04-21 13:59 /user/data/empdata_pig.csv
vboxuser@vanditha:~/hadoopdata$
```

Inserting data into external table student1 from the HDFS

```
Time taken: 0.076 seconds, Fetched: 4 row(s)
hive> load data inpath "/user/data/data.csv" into table student1;
Loading data to table db.student1
OK
Time taken: 0.219 seconds
hive>
```

```
Time taken: 0.219 seconds
hive> select * from student1;
OK
student1.sid      student1.sname  student1.marks  student1.djno
1                raj            80             12
2                ram            81             11
3                sam            88             12
4                rita            70             10
NULL            NULL           NULL           NULL
Time taken: 0.149 seconds, Fetched: 5 row(s)
hive>
```

Partion:

```
Time taken: 0.149 seconds, Fetched: 5 row(s)
hive> create external table if not exists ext_student(sid int,sname string,marks int) partitioned by (dpno int) row format delimited fields terminated by ',';
OK
Time tselect * from ext_student;
```

```
Time tselect * from ext_student;
OKve>
ext_student.sid  ext_student.sname      ext_student.marks      ext_student.djno
```

```
hive> set hive.exec.dynamic.partition.mode;
hive.exec.dynamic.partition.mode=strict
hive> set hive.exec.dynamic.partition.mode=unstrict;
hive>
```

Insert data to the table from existing one

```
hive> insert into table ext_student partition(dpno) select * from student;
Query ID = vboxuser_20230429001356_01a5c84a-6dd8-4b7f-92cb-0372a3334539
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
```